

Data Science for Python Folks Without (or With!) a Ph.D.

Douglas Starnes

Granite State Code Camp 2023



Spoiler Alert

You need to remember only three ideas today

- 1) Data is worthless
- 2) Analytics are worth pennies
- 3) Decisions are worth dollars

Obligatory Narcissism Slide (who is this guy?)

- Hi! I'm Douglas!
- Memphis, TN area
- Entrepreneur and technical author
- 4x Microsoft MVP (DevTech/Python)
- Memphis Python & Memphis Azure
- Scenic City Summit
- TDevConf
- Trained composer, board game and LEGO collector

Commercial Break

- Memphis Azure User Group - virtual
 - Monday, Dec. 4 at 6:00 PM Central
Matt Eland – Automating My Dog with Azure AI Services
RSVP: <https://bit.ly/memazug-dec-23>
- Memphis Python User Group – virtual
 - Monday, Dec. 18 at 6:00 PM Central
New Year's Resolutions – Launch That Startup
RSVP: <https://bit.ly/mempy-dec-23>
- Interested in speaking?
 - douglasastarnes@outlook.com
 - <https://linktr.ee/douglasstarnes> (LinkedIn is best)
- DataTune 2024
 - March 8-9, 2024, Nashville, TN
<https://datatuneconf.com>

Agenda

Python

numpy

pandas

matplotlib

Jupyter
Notebook

Python

Why Python?

- According to GitHub, it is the 2nd most used language
- Simple, clean syntax
- Easy to learn, read, and remember
- Open source
- Cross platform
- Numerous applications
- Large number of 3rd party packages

Slicing strings

```
String sample = "Granite State Code Camp"

// "Granite" (first 7)
System.out.println(sample.Substring(0, 7));

// "State"
System.out.println(sample.Substring(8, 13));

// "Camp" (last 4)
System.out.println(sample.Substring(sample.length() - 4));
```

Java

```
sample = "Granite State Code Camp"

# "Granite" (first 7)
print(sample[:7])

# "State"
print(sample[8:13])

# "Camp" (last 4)
print(sample[-4:])
```

Python

Reversing a string

```
String sample = "Granite State Code Camp";
char[] charArray = sample.toCharArray();
int left = 0;
int right = charArray.length - 1;

while (left < right) {
    char temp = charArray[left];
    charArray[left] = charArray[right];
    charArray[right] = temp;

    left++;
    right--;
}
```

Java

```
sample = "Granite State Code Camp"

reversed = ""

for c in sample:
    reversed = c + reversed

print(reversed)
```

Python

Reversing a string

```
String sample = "Granite State Code Camp";
char[] charArray = sample.toCharArray();
int left = 0;
int right = charArray.length - 1;

while (left < right) {
    char temp = charArray[left];
    charArray[left] = charArray[right];
    charArray[right] = temp;

    left++;
    right--;
}
```

Java

```
sample = "Granite State Code Camp"

reversed = list(sample)

reversed.reverse()

print("".join(reversed))
```

Python

Reversing a string

```
String sample = "Granite State Code Camp";  
char[] charArray = sample.toCharArray();  
int left = 0;  
int right = charArray.length - 1;  
  
while (left < right) {  
    char temp = charArray[left];  
    charArray[left] = charArray[right];  
    charArray[right] = temp;  
  
    left++;  
    right--;  
}
```

Java

```
sample = "Granite State Code Camp"  
  
print(sample[::-1])
```

Python

Methods / Functions

```
public static int[] divide(int a, int b) {  
    int[] result = new int[2];  
    result[0] = a / b;  
    result[1] = a % b;  
    return result;  
}  
  
public static void Main() {  
    int[] result = divide(5, 3);  
    int quotient = result[0];  
    int remainder = result[1];  
    // display  
}
```

Java

```
def divide(a, b):  
    return (a // b, a % b)  
  
result = divide(5, 3)  
  
quotient = result[0]  
remainder = result[1]  
  
print(f"q: {quotient}, r: {remainder}")
```

Python

Methods / Functions

```
public static int[] divide(int a, int b) {  
    int[] result = new int[2];  
    result[0] = a / b;  
    result[1] = a % b;  
    return result;  
}  
  
public static void Main() {  
    int[] result = divide(5, 3);  
    int quotient = result[0];  
    int remainder = result[1];  
    // display  
}
```

Java

```
def divide(a, b):  
    return (a // b, a % b)  
  
quotient, remainder = divide(5, 3)  
  
print(f"q: {quotient}, r: {remainder}")
```

Python

Methods / Functions

```
public static int[] divide(int a, int b) {  
    int[] result = new int[2];  
    result[0] = a / b;  
    result[1] = a % b;  
    return result;  
}  
  
public static void Main() {  
    int[] result = divide(5, 3);  
    int quotient = result[0];  
    int remainder = result[1];  
    // display  
}
```

Java

```
from collections import namedtuple  
  
Result = namedtuple("Result", ["quotient", "remainder"])  
  
def divide(a, b):  
    return Result(a // b, a % b)  
  
result = divide(5, 3)  
  
print(f"q: {result.quotient}, r: {result.remainder}")
```

Python

Methods / Functions

```
public static int[] divide(int a, int b) {  
    int[] result = new int[2];  
    result[0] = a / b;  
    result[1] = a % b;  
    return result;  
}  
  
public static void Main() {  
    int[] result = divide(5, 3);  
    int quotient = result[0];  
    int remainder = result[1];  
    // display  
}
```

Java

```
from collections import namedtuple  
  
Result = namedtuple("Result", ["quotient", "remainder"])  
  
def divide(a, b):  
    return Result(a // b, a % b)  
  
try:  
    result = divide(5, 0)  
    print(f"q: {result.quotient}, r: {result.remainder}")  
except ZeroDivisionError as e:  
    print(e)
```

Python

The Zen of Python, by Tim Peters

Beautiful is better than ugly.

Explicit is better than implicit.

Simple is better than complex.

Complex is better than complicated.

Flat is better than nested.

Sparse is better than dense.

Readability counts.

Special cases aren't special enough to break the rules.

Although practicality beats purity.

Errors should never pass silently.

Unless explicitly silenced.

In the face of ambiguity, refuse the temptation to guess.

There should be one-- and preferably only one --obvious way to do it.

Although that way may not be obvious at first unless you're Dutch.

Now is better than never.

Although never is often better than *right* now.

If the implementation is hard to explain, it's a bad idea.

If the implementation is easy to explain, it may be a good idea.

Namespaces are one honking great idea -- let's do more of those!

The Zen of Python, by Tim Peters

Beautiful is better than ugly.

Explicit is better than implicit.

Simple is better than complex.

Complex is better than complicated.

Flat is better than nested.

Sparse is better than dense.

Readability counts.

Special cases aren't special enough to break the rules.

Although practicality beats purity.

Errors should never pass silently.

Unless explicitly silenced.

In the face of ambiguity, refuse the temptation to guess.

There should be one-- and preferably only one --obvious way to do it.

Although that way may not be obvious at first unless you're Dutch.

Now is better than never.

Although never is often better than *right* now.

If the implementation is hard to explain, it's a bad idea.

If the implementation is easy to explain, it may be a good idea.

Namespaces are one honking great idea -- let's do more of those!

The Zen of Python, by Tim Peters

Beautiful is better than ugly.

Explicit is better than implicit.

Simple is better than complex.

Complex is better than complicated.

Flat is better than nested.

Sparse is better than dense.

Readability counts.

Special cases aren't special enough to break the rules.

Although practicality beats purity.

Errors should never pass silently.

Unless explicitly silenced.

In the face of ambiguity, refuse the temptation to guess.

There should be one-- and preferably only one --obvious way to do it.

Although that way may not be obvious at first unless you're Dutch.

Now is better than never.

Although never is often better than *right* now.

If the implementation is hard to explain, it's a bad idea.

If the implementation is easy to explain, it may be a good idea.

Namespaces are one honking great idea -- let's do more of those!

The Zen of Python, by Tim Peters

Beautiful is better than ugly.

Explicit is better than implicit.

Simple is better than complex.

Complex is better than complicated.

Flat is better than nested.

Sparse is better than dense.

Readability counts.

Special cases aren't special enough to break the rules.

Although practicality beats purity.

Errors should never pass silently.

Unless explicitly silenced.

In the face of ambiguity, refuse the temptation to guess.

There should be one-- and preferably only one --obvious way to do it.

Although that way may not be obvious at first unless you're Dutch.

Now is better than never.

Although never is often better than *right* now.

If the implementation is hard to explain, it's a bad idea.

If the implementation is easy to explain, it may be a good idea.

Namespaces are one honking great idea -- let's do more of those!

The Zen of Python, by Tim Peters

Beautiful is better than ugly.

Explicit is better than implicit.

Simple is better than complex.

Complex is better than complicated.

Flat is better than nested.

Sparse is better than dense.

Readability counts.

Special cases aren't special enough to break the rules.

Although practicality beats purity.

Errors should never pass silently.

Unless explicitly silenced.

In the face of ambiguity, refuse the temptation to guess.

There should be one-- and preferably only one --obvious way to do it.

Although that way may not be obvious at first unless you're Dutch.

Now is better than never.

Although never is often better than *right* now.

If the implementation is hard to explain, it's a bad idea.

If the implementation is easy to explain, it may be a good idea.

Namespaces are one honking great idea -- let's do more of those!

The Zen of Python, by Tim Peters

Beautiful is better than ugly.

Explicit is better than implicit.

Simple is better than complex.

Complex is better than complicated.

Flat is better than nested.

Sparse is better than dense.

Readability counts.

Special cases aren't special enough to break the rules.

Although practicality beats purity.

Errors should never pass silently.

Unless explicitly silenced.

In the face of ambiguity, refuse the temptation to guess.

There should be one-- and preferably only one --obvious way to do it.

Although that way may not be obvious at first unless you're Dutch.

Now is better than never.

Although never is often better than *right* now.

If the implementation is hard to explain, it's a bad idea.

If the implementation is easy to explain, it may be a good idea.

Namespaces are one honking great idea -- let's do more of those!

numpy

1) Data is worthless

```
nums = [2, 4, 8, 7, 3,  
nums_inc = nums + 1
```

TypeError

Python doesn't know if you
want to append 1 to the list
or add 1 to each element in
the list

```
nums = [2, 4, 8, 7, 3, 1]
nums_inc = [num + 1 for num in nums]
nums_inc = nums + [1]
import numpy as np
num_arr = np.array(nums)
num_arr_inc = num_arr + 1
```



```
array([3, 5, 9, 8, 4, 2])
```

The fundamental data structure in numpy is the ndarray.

A numpy ndarray is like a Python `list` that has superpowers.

A numpy ndarray can have multiple dimensions.

The fundamental data structure in numpy is the ndarray.

A numpy ndarray is like a Python list that has superpowers.

A numpy ndarray can have multiple dimensions.

```
import random
```

```
def random_matrix(rows=5, cols=4):
```

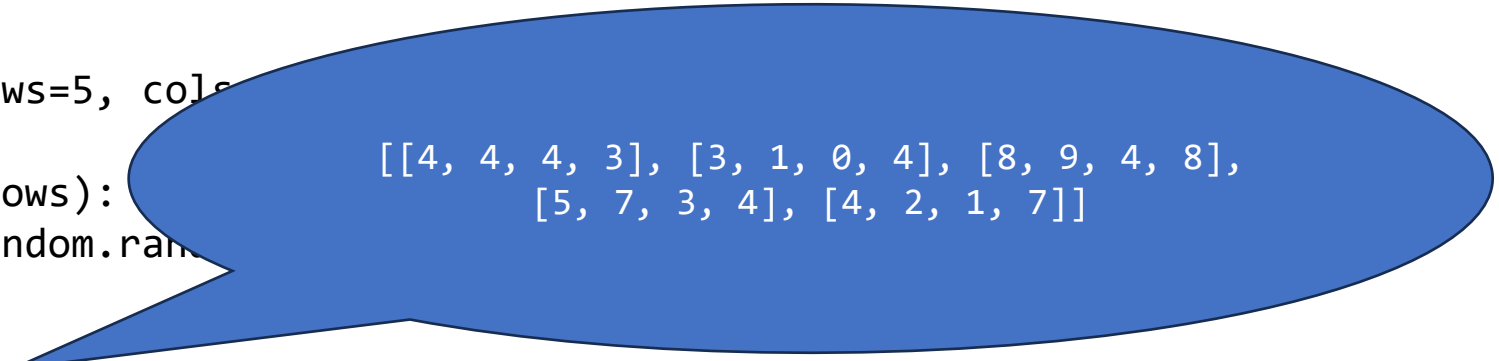
```
    A = []
```

```
    for _ in range(rows):
```

```
        A.append([random.random() for _ in range(cols)])
```

```
    return A
```

```
M = random_matrix()
```



```
[[4, 4, 4, 3], [3, 1, 0, 4], [8, 9, 4, 8],  
 [5, 7, 3, 4], [4, 2, 1, 7]]
```

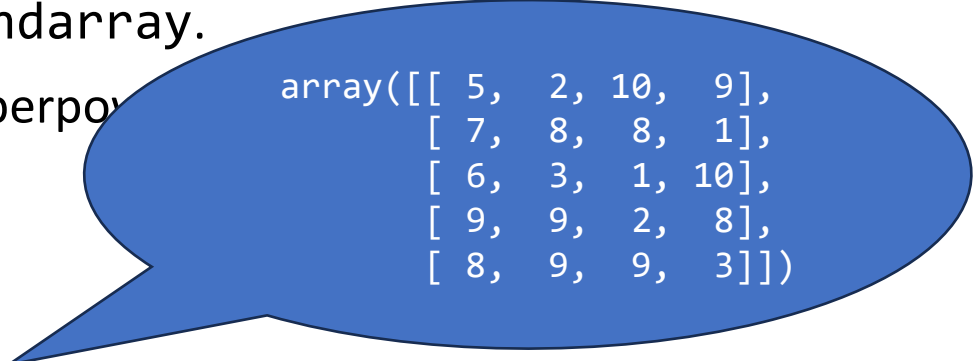
The fundamental data structure in numpy is the ndarray.

A numpy ndarray is like a Python list that has superpowers.

A numpy ndarray can have multiple dimensions.

```
import numpy as np
```

```
M = np.random.randint(0, 11, (5, 4))
```



```
array([[ 5,  2, 10,  9],  
       [ 7,  8,  8,  1],  
       [ 6,  3,  1, 10],  
       [ 9,  9,  2,  8],  
       [ 8,  9,  9,  3]])
```

The fundamental data structure in numpy is the ndarray.

A numpy ndarray is like a Python list that has superpowers.

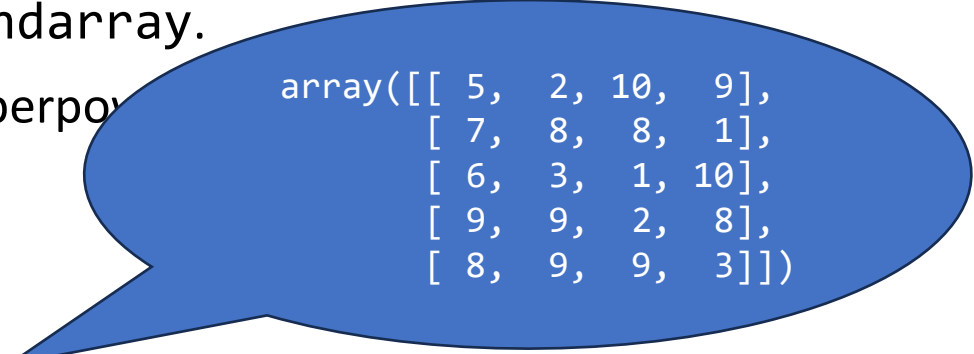
A numpy ndarray can have multiple dimensions.

```
import numpy as np
```

```
M = np.random.randint(0, 11, (5, 4))
```

You can access individual elements and slice an ndarray using list syntax:

```
M[1][3]
```



```
array([[ 5,  2, 10,  9],  
       [ 7,  8,  8,  1],  
       [ 6,  3,  1, 10],  
       [ 9,  9,  2,  8],  
       [ 8,  9,  9,  3]])
```



1

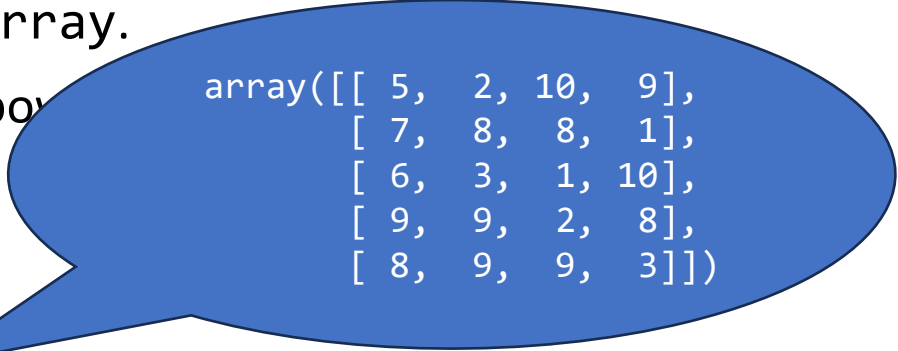
The fundamental data structure in numpy is the ndarray.

A numpy ndarray is like a Python list that has superpowers.

A numpy ndarray can have multiple dimensions.

```
import numpy as np
```

```
M = np.random.randint(0, 11, (5, 4))
```



```
array([[ 5,  2, 10,  9],  
       [ 7,  8,  8,  1],  
       [ 6,  3,  1, 10],  
       [ 9,  9,  2,  8],  
       [ 8,  9,  9,  3]])
```

You can access individual elements and slice an ndarray using list syntax:

```
M[1][3]
```

An ndarray has a special and cleaner syntax for elements in multiple dimensions

```
M[1, 3]
```



1

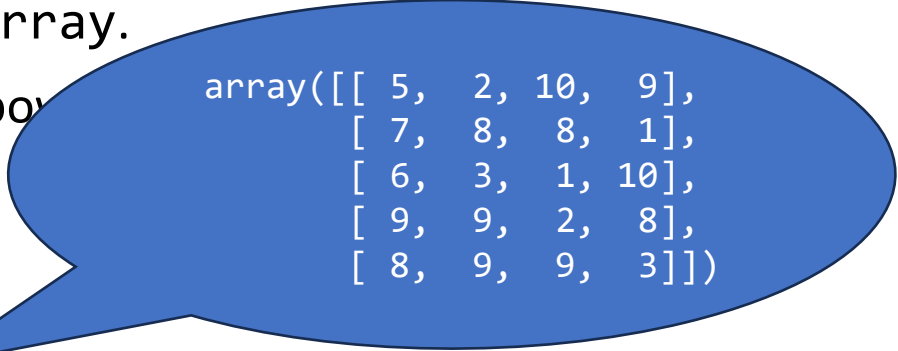
The fundamental data structure in numpy is the ndarray.

A numpy ndarray is like a Python list that has superpowers.

A numpy ndarray can have multiple dimensions.

```
import numpy as np
```

```
M = np.random.randint(0, 11, (5, 4))
```



```
array([[ 5,  2, 10,  9],  
       [ 7,  8,  8,  1],  
       [ 6,  3,  1, 10],  
       [ 9,  9,  2,  8],  
       [ 8,  9,  9,  3]])
```

You can access individual elements and slice an ndarray using list syntax

```
M[1][3]
```

An ndarray has a special and cleaner syntax for elements in multiple dimensions

```
M[1, 3]
```

You can also use slice notation to access elements of an ndarray

```
M[:, 3]
```



```
array([ 9,  1, 10,  8,  3])
```

An ndarray is aware of

M.shape



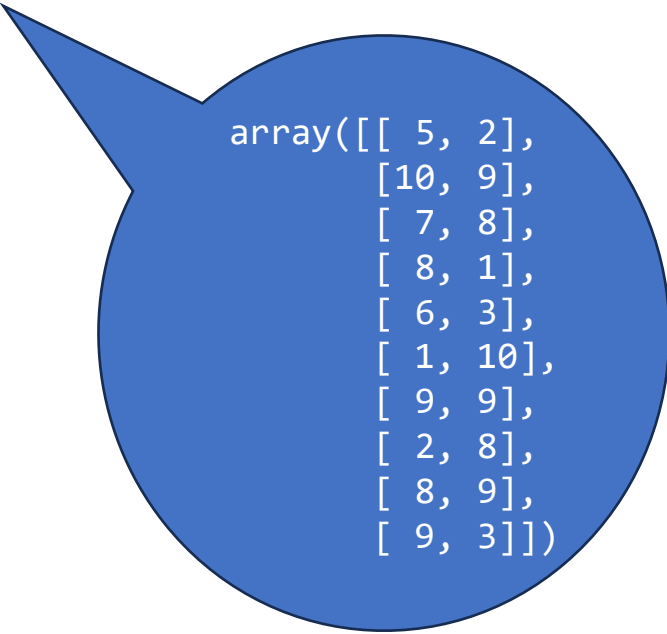
(5, 4)

An ndarray is aware of its shape

`M.shape`

The shape of an ndarray can be modified (the product of the new and old dimensions must be equal)

`M.reshape(10, 2)`



```
array([[ 5, 2],  
       [10, 9],  
       [ 7, 8],  
       [ 8, 1],  
       [ 6, 3],  
       [ 1, 10],  
       [ 9, 9],  
       [ 2, 8],  
       [ 8, 9],  
       [ 9, 3]])
```

An ndarray is aware of its shape

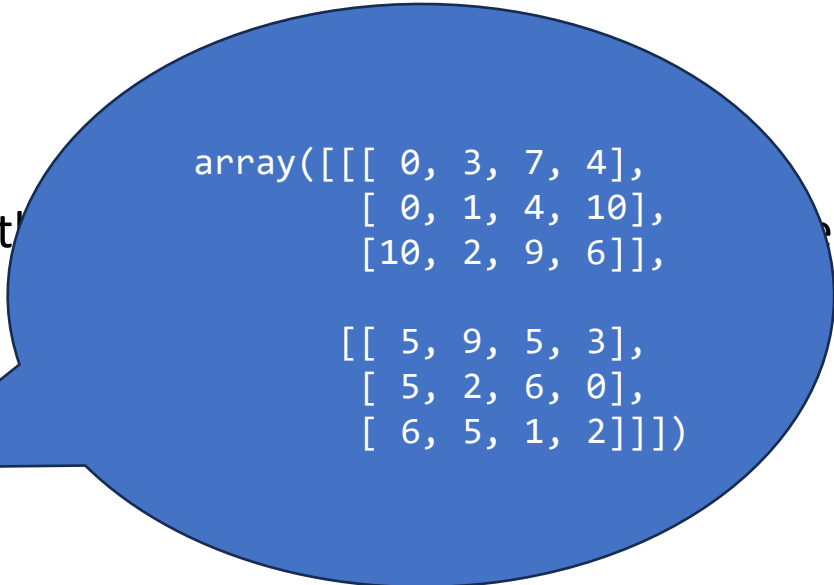
```
M.shape
```

The shape of an ndarray can be modified (the product of the

```
M.reshape(10, 2)
```

The features of an ndarray scale to infinite dimensions

```
M = np.random.randint(0, 11, (2, 3, 4))
```



```
array([[[ 0, 3, 7, 4],  
       [ 0, 1, 4, 10],  
       [10, 2, 9, 6]],  
      [[ 5, 9, 5, 3],  
       [ 5, 2, 6, 0],  
       [ 6, 5, 1, 2]]])
```

equal)

An ndarray is aware of its shape

```
M.shape
```

The shape of an ndarray can be modified (the product of the

```
M.reshape(10, 2)
```

The features of an ndarray scale to infinite dimensions

```
M = np.random.randint(0, 11, (2, 3, 4))
```

```
M[1, 0, 3]
```

3

```
array([[[ 0, 3, 7, 4],  
       [ 0, 1, 4, 10],  
       [10, 2, 9, 6]],  
      [[ 5, 9, 5, 3],  
       [ 5, 2, 6, 0],  
       [ 6, 5, 1, 2]]])
```

An ndarray is aware of its shape

```
M.shape
```

The shape of an ndarray can be modified (the product of the original shape must be equal to the new shape)

```
M.reshape(10, 2)
```

The features of an ndarray scale to infinite dimensions

```
M = np.random.randint(0, 11, (2, 3, 4))
```

```
M[1, 0, 3]
```

```
M[:, 0, :]
```

```
array([[0, 3, 7, 4],  
       [5, 9, 5, 3]])
```

```
array([[[ 0, 3, 7, 4],  
        [ 0, 1, 4, 10],  
        [10, 2, 9, 6]],  
       [[ 5, 9, 5, 3],  
        [ 5, 2, 6, 0],  
        [ 6, 5, 1, 2]]])
```


An ndarray is aware of its shape

```
M.shape
```

The shape of an ndarray can be modified (the product of the original shape must be equal to the new shape)

```
M.reshape(10, 2)
```

The features of an ndarray scale to infinite dimensions

```
M = np.random.randint(0, 11, (2, 3, 4))
```

```
M[1, 0, 3]
```

```
M[:, 0, :]
```

```
M.reshape(4, 2, 3)
```

```
array([[[ 0, 3, 7, 4],
        [ 0, 1, 4, 10],
        [10, 2, 9, 6]],
       [[ 5, 9, 5, 3],
        [ 5, 2, 6, 0],
        [ 6, 5, 1, 2]]])
```

```
array([[[ 0, 3, 7],
        [ 4, 0, 1]],
       [[ 4, 10, 10],
        [ 2, 9, 6]],
       [[ 5, 9, 5],
        [ 3, 5, 2]],
       [[ 6, 0, 6],
        [ 5, 1, 2]]])
```

An ndarray is aware of its shape

```
M.shape
```

The shape of an ndarray can be modified (the product of the

```
M.reshape(10, 2)
```

The features of an ndarray scale to infinite dimensions

```
M = np.random.randint(0, 11, (2, 3, 4))
```

```
M[1, 0, 3]
```

```
M[:, 0, :]
```

```
M.reshape(4, 2, 3)
```

```
M.reshape(2, 2, 2, 3)
```

```
array([[[ 0, 3, 7, 4],
        [ 0, 1, 4, 10],
        [10, 2, 9, 6]],
       [[ 5, 9, 5, 3],
        [ 5, 2, 6, 0],
        [ 6, 5, 1, 2]]])
```

```
array([[[[ 0, 3, 7],
         [ 4, 0, 1]],
        [[ 4, 10, 10],
         [ 2, 9, 6]]],
       [[[ 5, 9, 5],
         [ 3, 5, 2]],
        [[ 6, 0, 6],
         [ 5, 1, 2]]]])
```

An ndarray is aware of its shape

```
M.shape
```

The shape of an ndarray can be modified (the product of the

```
M.reshape(10, 2)
```

The features of an ndarray scale to infinite dimensions

```
M = np.random.randint(0, 11, (2, 3, 4))
```

```
M[1, 0, 3]
```

```
M[:, 0, :]
```

```
M.reshape(4, 2, 3)
```

```
M.reshape(2, 2, 2, 3)
```

```
M.reshape(8, 3)
```

```
array([[[ 0, 3, 7, 4],
        [ 0, 1, 4, 10],
        [10, 2, 9, 6]],
       [[ 5, 9, 5, 3],
        [ 5, 2, 6, 0],
        [ 6, 5, 1, 2]]])
```

```
array([[ 0, 3, 7],
       [ 4, 0, 1],
       [ 4, 10, 10],
       [ 2, 9, 6],
       [ 5, 9, 5],
       [ 3, 5, 2],
       [ 6, 0, 6],
       [ 5, 1, 2]])
```

An ndarray is aware of its shape

```
M.shape
```

The shape of an ndarray can be modified (the product of the original shape must be equal to the new shape)

```
M.reshape(10, 2)
```

The features of an ndarray scale to infinite dimensions

```
M = np.random.randint(0, 11, (2, 3, 4))
```

```
M[1, 0, 3]
```

```
M[:, 0, :]
```

```
M.reshape(4, 2, 3)
```

```
M.reshape(2, 2, 2, 3)
```

```
M.reshape(8, 3)
```

```
M.flatten()
```

```
array([[[ 0, 3, 7, 4],  
        [ 0, 1, 4, 10],  
        [10, 2, 9, 6]],  
       [[ 5, 9, 5, 3],  
        [ 5, 2, 6, 0],  
        [ 6, 5, 1, 2]]])
```

```
array([ 0, 3, 7, 4, 0, 1, 4, 10, 10, 2, 9,  
       6, 5, 9, 5, 3, 5, 2, 6, 0, 6, 5, 1, 2])
```

pandas

2) Analytics are worth pennies

The fundamental data structure in pandas is the DataFrame.

You can create a DataFrame from many sources.

```
import pandas as pd
```

```
games_df = pd.read_csv("games.csv")
```

```
year,week,home,away,win,loss,pts_win,pts_loss
2000,1,Vikings,Bears,Vikings,Bears,30,27
2000,1,Chiefs,Colts,Colts,Chiefs,27,14
2000,1,Redskins,Panthers,Redskins,Panthers,20,17
2000,1,Falcons,49ers,Falcons,49ers,36,28
2000,1,Steelers,Ravens,Ravens,Steelers,16,0
2000,1,Browns,Jaguars,Jaguars,Browns,27,7
2000,1,Patriots,Buccaneers,Buccaneers,Patriots,21,16
```

- read_csv()
- read_excel()
- read_html()
- read_json()
- read_pickle()
- read_sql()

The fundamental data structure in pandas is the DataFrame.

You can create a DataFrame from many sources.

```
import pandas as pd
```

```
games_df = pd.read_csv('games.csv')
```

```
games_df.head()
```

	year	week	home	away	win	loss	pts_win	pts_loss
0	2000	1	Vikings	Bears	Vikings	Bears	30	27
1	2000	1	Chiefs	Colts	Colts	Chiefs	27	14
2	2000	1	Redskins	Panthers	Redskins	Panthers	20	17
3	2000	1	Falcons	49ers	Falcons	49ers	36	28
4	2000	1	Steelers	Ravens	Ravens	Steelers	16	0

The fundamental data structure in

You can create a DataFrame from

```
import pandas as pd
```

```
games_df = pd.read_csv('nba_games.csv')
```

```
games_df.head()
```

```
games_df.info()
```

```
games_df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 5324 entries, 0 to 5323  
Data columns (total 8 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   year        5324 non-null   int64  
1   week        5324 non-null   int64  
2   home        5324 non-null   object  
3   away        5324 non-null   object  
4   win         5324 non-null   object  
5   loss        5324 non-null   object  
6   pts_win     5324 non-null   int64  
7   pts_loss    5324 non-null   int64  
dtypes: int64(4), object(4)  
memory usage: 332.9+ KB
```

	year	week	pts_win	pts_loss
count	5324.000000	5324.000000	5324.000000	5324.000000
mean	2009.527047	9.511833	27.781555	16.088843
std	5.754236	5.271909	8.830090	8.137451
min	2000.000000	1.000000	3.000000	0.000000
25%	2005.000000	5.000000	21.000000	10.000000
50%	2010.000000	10.000000	27.000000	16.000000
75%	2015.000000	14.000000	34.000000	21.000000
max	2019.000000	21.000000	62.000000	51.000000

The fundamental data structure in pandas is the DataFrame.

You can create a DataFrame from many sources.

```
import pandas as pd
```

```
games_df = pd.read_csv("games.csv")
```

```
games_df.head()
```

```
games_df.info()
```

```
games_df.describe()
```

Individual columns of a DataFrame

```
games_df["home"]
```

```
0      Vikings
1      Chiefs
2    Redskins
3      Falcons
4    Steelers
...
5319   Chiefs
5320  Packers
5321   Chiefs
5322    49ers
5323   Chiefs
Name: home, Length: 5324, dtype: object
```

The fundamental data structure in pandas is the DataFrame.

You can create a DataFrame from many sources.

```
import pandas as pd
```

```
games_df = pd.read_csv("games.csv")
```

```
games_df.head()
```

```
games_df.info()
```

```
games_df.describe()
```

Individual columns of a DataFrame

```
games_df["home"]
```

```
0    Vikings
1    Chiefs
2    Redskins
3    Falcons
4    Steelers
...
5319  Chiefs
5320  Packers
5321  Chiefs
5322   49ers
5323  Chiefs
Name: home, Length: 5324, dtype: object
```

The data structure for a column is a `Series`, which is like a Python `list` with an index

A `DataFrame` is a collection of `Series` that share an index

You can access individual rows by index:

```
games_df.loc[1]
```

Or by zero based position

```
games_df.iloc[1]
```

In this case they are the same.

A row is a Series

```
year      2000
week       1
home      Chiefs
away      Colts
win        Colts
loss      Chiefs
pts_win    27
pts_loss   14
Name: 1, dtype: object
```

You can access individual rows by index:

```
games_df.loc[1]
```

Or by zero based position

```
games_df.iloc[1]
```

In this case they are the same.

A row is a Series

You can also slice the rows, like a Python list (returns a DataFrame)

```
games_df.iloc[1:5]
```

	year	week	home	away	win	loss	pts_win	pts_loss
1	2000	1	Chiefs	Colts	Colts	Chiefs	27	14
2	2000	1	Redskins	Panthers	Redskins	Panthers	20	17
3	2000	1	Falcons	49ers	Falcons	49ers	36	28
4	2000	1	Steelers	Ravens	Ravens	Steelers	16	0

You can filter a column with a Boolean expression:

```
titans_home_games = games_df["home"] == "Titans"  
titans_home_games[0], titans_home_games[22]
```

(False, True)

	year	week	home	away	win	loss	pts_win	pts_loss
0	2000	1	Vikings	Bears	Vikings	Bears	30	27
22	2000	2	Titans	Chiefs	Titans	Chiefs	17	14

You can filter a column with a Boolean expression:

```
titans_home_games = games_df["home"] == "Titans"  
titans_home_games[0], titans_home_games[22]  
games_df[titans_home_games]
```

	year	week	home	away	win	loss	pts_win	pts_loss
22	2000	2	Titans	Chiefs	Titans	Chiefs	17	14
62	2000	5	Titans	Giants	Titans	Giants	28	14
99	2000	7	Titans	Jaguars	Titans	Jaguars	27	13
134	2000	10	Titans	Steelers	Titans	Steelers	9	7
150	2000	11	Titans	Ravens	Ravens	Titans	24	23
...
5169	2019	8	Titans	Buccaneers	Titans	Buccaneers	27	23
5197	2019	10	Titans	Chiefs	Titans	Chiefs	35	32
5229	2019	12	Titans	Jaguars	Titans	Jaguars	42	20
5272	2019	15	Titans	Texans	Texans	Titans	24	21
5287	2019	16	Titans	Saints	Saints	Titans	38	28

[163 rows x 8 columns]

You can filter a column with a Boolean expression:

```
titans_home_games = games_df["home"] == "Titans"
titans_home_games[0], titans_home_games[22]
games_df[titans_home_games]
```

Use the and (&) and or (|) operators to combine filters:

```
titans_wins = games_df["win"] == "Titans"
titans_home_wins = games_df[titans_home_games & titans_wins]
```

	year	week	home	away	win	loss	pts_win	pts_loss
22	2000	2	Titans	Chiefs	Titans	Chiefs	17	14
62	2000	5	Titans	Giants	Titans	Giants	28	14
99	2000	7	Titans	Jaguars	Titans	Jaguars	27	13
134	2000	10	Titans	Steelers	Titans	Steelers	9	7
165	2000	12	Titans	Browns	Titans	Browns	24	10
...
5014	2018	16	Titans	Redskins	Titans	Redskins	25	16
5158	2019	7	Titans	Chargers	Titans	Chargers	23	20
5169	2019	8	Titans	Buccaneers	Titans	Buccaneers	27	23
5197	2019	10	Titans	Chiefs	Titans	Chiefs	35	32
5229	2019	12	Titans	Jaguars	Titans	Jaguars	42	20

[88 rows x 8 columns]

You can filter a column with a Boolean expression:

```
titans_home_games = games_df["home"] == "Titans"  
titans_home_games[0], titans_home_games[22]  
games_df[titans_home_games]
```

Use the and (&) and or (|) operators to combine filters:

```
titans_wins = games_df["win"] == "Titans"  
titans_home_wins = games_df[titans_home_games & titans_wins]
```

Compute statistics on a column:

```
titans_home_wins["pts_win"].mean()
```

26.806818181818183

You can filter a column with a Boolean expression:

```
titans_home_games = games_df["home"] == "Titans"  
titans_home_games[0], titans_home_games[22]  
games_df[titans_home_games]
```

Use the and (&) and or (|) operators to combine filters:

```
titans_wins = games_df["win"] == "Titans"  
titans_home_wins = games_df[titans_home_games & titans_wins]
```

Compute statistics on a column:

```
titans_home_wins["pts_win"].mean()
```

Add a new column:

```
titans_home_wins["spread"] = titans_home_wins["pts_win"] - titans_home_wins["pts_loss"]  
titans_home_wins.head()
```

	year	week	home	away	win	loss	pts_win	pts_loss	spread
22	2000	2	Titans	Chiefs	Titans	Chiefs	17	14	3
62	2000	5	Titans	Giants	Titans	Giants	28	14	14
99	2000	7	Titans	Jaguars	Titans	Jaguars	27	13	14
134	2000	10	Titans	Steelers	Titans	Steelers	9	7	2
165	2000	12	Titans	Browns	Titans	Browns	24	10	14

You can filter a column with a Boolean expression:

```
titans_home_games = games_df["home"] == "Titans"
titans_home_games[0], titans_home_games[22]
games_df[titans_home_games]
```

Use the and (&) and or (|) operators to combine filters:

```
titans_wins = games_df["win"] == "Titans"
titans_home_wins = games_df[titans_home_games & titans_wins]
```

Compute statistics on a column:

```
titans_home_wins["pts_win"].mean()
```

Add a new column:

```
titans_home_wins["spread"] = titans_home_wins["pts_win"] - titans_home_wins["pts_loss"]
titans_home_wins.head()
```

Sort by spread descending:

```
titans_home_wins.sort_values("spread", ascending=False).head()
```

	year	week	home	away	win	loss	pts_win	pts_loss	spread
2590	2009	14	Titans	Rams	Titans	Rams	47	7	40
210	2000	15	Titans	Bengals	Titans	Bengals	35	3	32
247	2000	17	Titans	Cowboys	Titans	Cowboys	31	0	31
2661	2010	1	Titans	Raiders	Titans	Raiders	38	13	25
3512	2013	4	Titans	Jets	Titans	Jets	38	13	25

You can group a DataFrame by one or more columns

```
pts_grps = games_df.groupby(["year", "win"])
```

And compute summary statistics

```
pts_sum = pts_grps.sum()[["pts_win", "pts_loss"]]
```

The index has multiple values

```
pts_sum.loc[2019, "Titans"]
```

```
pts_win    350
pts_loss    195
Name: (2019, Titans),
dtype: int64
```

		pts_win	pts_loss
year	win		
2000	49ers	167	74
	Bears	114	95
	Bengals	84	51
	Bills	192	131
	Broncos	380	226
...	
2019	Seahawks	333	263
	Steelers	180	110
	Texans	294	214
	Titans	350	195
	Vikings	329	179

[636 rows x 2 columns]

You can group a DataFrame by one or more columns

```
pts_grps = games_df.groupby(["year", "win"])
```

And compute summary statistics

```
pts_sum = pts_grps.sum()[["pts_win", "pts_loss"]]
```

The index has multiple values

```
pts_sum.loc[2019, "Titans"]
```

The DataFrame can be sorted with the grouped columns

```
pts_sum.sort_values(["year", "win"], ascending=[True, False])
```

		pts_win	pts_loss
year	win		
2000	Raiders	414	191
	Ravens	407	136
	Rams	400	272
	Broncos	380	226
	Vikings	331	216
...	
2019	Dolphins	144	120
	Giants	133	89
	Lions	98	87
	Redskins	65	53
	Bengals	55	29

[636 rows x 2 columns]

matplotlib

3) Decisions are worth dollars

Let's look at the weekly attendance of Titans games in 2000

```
import pandas as pd
```

```
attend_df = pd.read_csv("attendance.csv")
```

```
titans_attendance = attend_df[attend_df["team_name"] == "Titans"]
```

	team	team_name	year	total	home	away	week	weekly_attendance
493	Tennessee	Titans	2000	1091274	547524	543750	1	72492.0
494	Tennessee	Titans	2000	1091274	547524	543750	2	68203.0
495	Tennessee	Titans	2000	1091274	547524	543750	3	NaN
496	Tennessee	Titans	2000	1091274	547524	543750	4	51769.0
497	Tennessee	Titans	2000	1091274	547524	543750	5	68341.0
...
10824	Tennessee	Titans	2019	1047496	516074	531422	13	60361.0
10825	Tennessee	Titans	2019	1047496	516074	531422	14	52760.0
10826	Tennessee	Titans	2019	1047496	516074	531422	15	65265.0
10827	Tennessee	Titans	2019	1047496	516074	531422	16	66756.0
10828	Tennessee	Titans	2019	1047496	516074	531422	17	71794.0

[340 rows x 8 columns]

Let's look at the weekly attendance of Titans games in 2000

```
import pandas as pd
```

```
attend_df = pd.read_csv("attendance.csv")
```

```
titans_attendance = attend_df[attend_df["team_name"] == "Titans"]
```

```
bye_weeks = titans_attendance["weekly_attendance"].isna()
```

```
titans_attendance.drop(titans_attendance[bye_weeks].index, inplace=True)
```

	team	team_name	year	total	home	away	week	weekly_attendance
493	Tennessee	Titans	2000	1091274	547524	543750	1	72492.0
494	Tennessee	Titans	2000	1091274	547524	543750	2	68203.0
496	Tennessee	Titans	2000	1091274	547524	543750	4	51769.0
497	Tennessee	Titans	2000	1091274	547524	543750	5	68341.0
498	Tennessee	Titans	2000	1091274	547524	543750	6	63406.0
...
10824	Tennessee	Titans	2019	1047496	516074	531422	13	60361.0
10825	Tennessee	Titans	2019	1047496	516074	531422	14	52760.0
10826	Tennessee	Titans	2019	1047496	516074	531422	15	65265.0
10827	Tennessee	Titans	2019	1047496	516074	531422	16	66756.0
10828	Tennessee	Titans	2019	1047496	516074	531422	17	71794.0

[320 rows x 8 columns]

Let's look at the weekly attendance of Titans games in 2000

```
import pandas as pd
```

```
att
```

```
ti
```

```
by
```

```
ti
```

```
ti
```

TIMES UP!


Which week in the 2000 season did the Titans have the highest attendance?

Let's look at the weekly attendance of Titans games in 2000

```
import pandas as pd
```

	team	team_name	year	total	home	away	week	weekly_attendance
493	Tennessee	Titans	2000	1091274	547524	543750	1	72492.0
494	Tennessee	Titans	2000	1091274	547524	543750	2	68203.0
496	Tennessee	Titans	2000	1091274	547524	543750	4	51769.0
497	Tennessee	Titans	2000	1091274	547524	543750	5	68341.0
498	Tennessee	Titans	2000	1091274	547524	543750	6	63406.0
499	Tennessee	Titans	2000	1091274	547524	543750	7	68498.0
500	Tennessee	Titans	2000	1091274	547524	543750	8	69200.0
501	Tennessee	Titans	2000	1091274	547524	543750	9	83472.0
502	Tennessee	Titans	2000	1091274	547524	543750	10	68498.0
503	Tennessee	Titans	2000	1091274	547524	543750	11	68490.0
504	Tennessee	Titans	2000	1091274	547524	543750	12	68498.0
505	Tennessee	Titans	2000	1091274	547524	543750	13	65454.0
506	Tennessee	Titans	2000	1091274	547524	543750	14	65639.0
507	Tennessee	Titans	2000	1091274	547524	543750	15	68498.0
508	Tennessee	Titans	2000	1091274	547524	543750	16	72318.0
509	Tennessee	Titans	2000	1091274	547524	543750	17	68498.0


Which week in the 2000 season did the Titans have the highest attendance?




Which week in the 2019 season did the Titans have the **highest** attendance?

Which week in the 2019 season did the Titans have the **lowest** attendance?


```
import matplotlib.pyplot as plt  
plt.bar(  
    np.arange(len(titans_attendance_2019)),  
    titans_attendance_2019["weekly_attendance"])
```




```
import matplotlib.pyplot as plt
plt.bar(
    np.arange(len(titans_attendance_2019)),
    titans_attendance_2019["weekly_attendance"])
plt.xticks(np.arange(len(titans_attendance_2019)), titans_attendance_2019["week"])
```




```
import matplotlib.pyplot as plt
plt.bar(
    np.arange(len(titans_attendance_2019)),
    titans_attendance_2019["weekly_attendance"])
plt.xticks(np.arange(len(titans_attendance_2019)), titans_attendance_2019["week"])
plt.legend(["Weekly Attendance"])
```




```
import matplotlib.pyplot as plt
plt.bar(
    np.arange(len(titans_attendance_2000)),
    titans_attendance_2000["weekly_attendance"])
plt.bar(
    np.arange(len(titans_attendance_2019)),
    titans_attendance_2019["weekly_attendance"])
plt.xticks(np.arange(len(titans_attendance_2019)), titans_attendance_2019["week"])
plt.legend(["Weekly Attendance"])
```




```
import matplotlib.pyplot as plt
plt.bar(
    np.arange(len(titans_attendance_2000)),
    titans_attendance_2000["weekly_attendance"], width=0.4)
plt.bar(
    np.arange(len(titans_attendance_2019)),
    titans_attendance_2019["weekly_attendance"], width=0.4)
plt.xticks(np.arange(len(titans_attendance_2019)), titans_attendance_2019["week"])
plt.legend(["Weekly Attendance"])
```



```
import matplotlib.pyplot as plt
plt.bar(
    np.arange(len(titans_attendance_2000)) - 0.2,
    titans_attendance_2000["weekly_attendance"], width=0.4)
plt.bar(
    np.arange(len(titans_attendance_2019)) + 0.2,
    titans_attendance_2019["weekly_attendance"], width=0.4)
plt.xticks(np.arange(len(titans_attendance_2019)), titans_attendance_2019["week"])
plt.legend(["Weekly Attendance"])
```




```
import matplotlib.pyplot as plt
plt.bar(
    np.arange(len(titans_attendance_2000)) - 0.2,
    titans_attendance_2000["weekly_attendance"], width=0.4)
plt.bar(
    np.arange(len(titans_attendance_2019)) + 0.2,
    titans_attendance_2019["weekly_attendance"], width=0.4)
plt.xticks(np.arange(len(titans_attendance_2019)), titans_attendance_2019["week"])
plt.legend(["2000", "2019"])
```




```
import matplotlib.pyplot as plt

plt.bar(
    np.arange(len(titans_attendance_2000)) - 0.2,
    titans_attendance_2000["weekly_attendance"], width=0.4)

plt.bar(
    np.arange(len(titans_attendance_2019)) + 0.2,
    titans_attendance_2019["weekly_attendance"], width=0.4)

plt.title("Titans Attendance 2000 vs 2019")
plt.xlabel("Week")
plt.ylabel("Attendance")

plt.xticks(np.arange(len(titans_attendance_2000) + 1))
plt.legend(["2000", "2019"])
```



```
import matplotlib.pyplot as plt

plt.bar(
    np.arange(len(titans_attendance_2000)) - 0.2,
    titans_attendance_2000["weekly_attendance"], width=0.4)

plt.bar(
    np.arange(len(titans_attendance_2019)) + 0.2,
    titans_attendance_2019["weekly_attendance"], width=0.4)

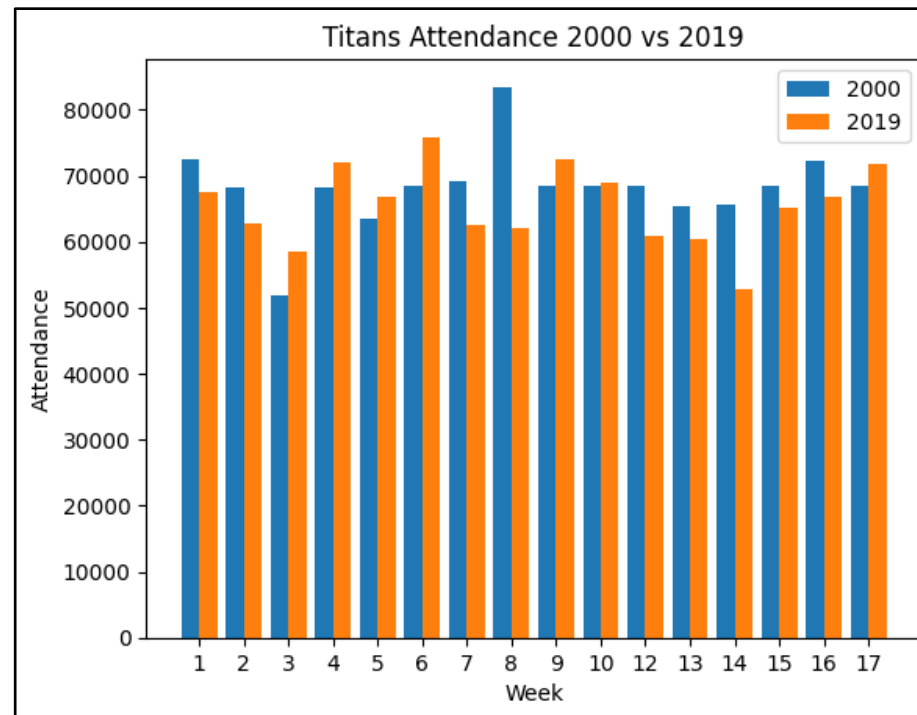
plt.title("Titans Attendance 2000 vs 2019")
plt.xlabel("Week")
plt.ylabel("Attendance")

plt.xticks(np.arange(len(titans_attendance_2019)), titans_attendance_2019["week"])

plt.legend(["2000", "2019"])
```

```
titans_no_game = attend_df[  
    (attend_df["team_name"] == "Titans") & (attend_df["weekly_attendance"].isna())]  
titans_no_game[titans_no_game["year"].isin([2000, 2019])]
```

	team	team_name	year	total	home	away	week	weekly_attendance
495	Tennessee	Titans	2000	1091274	547524	543750	3	NaN
10822	Tennessee	Titans	2019	1047496	516074	531422	11	NaN




Let's get the home and away attendance for each team in 2000

```
home_away_2000 = attend_df[(attend_df["year"] == 2000) & (attend_df["week"] == 1)]
```

	team	team_name	year	total	home	away	week	weekly_attendance
0	Arizona	Cardinals	2000	893926	387475	506451	1	77434.0
17	Atlanta	Falcons	2000	964579	422814	541765	1	54626.0
34	Baltimore	Ravens	2000	1062373	551695	510678	1	55049.0
51	Buffalo	Bills	2000	1098587	560695	537892	1	72492.0
68	Carolina	Panthers	2000	1095192	583489	511703	1	80257.0


Let's get the home and away attendance for each team in 2000

```
home_away_2000 = attend_df[(attend_df["year"] == 2000) & (attend_df["week"] == 1)]  
cardinals_2000 = home_away_2000[home_away_2000["team_name"] == "Cardinals"].iloc[0]  
plt.pie([cardinals_2000["home"], cardinals_2000["away"]])
```



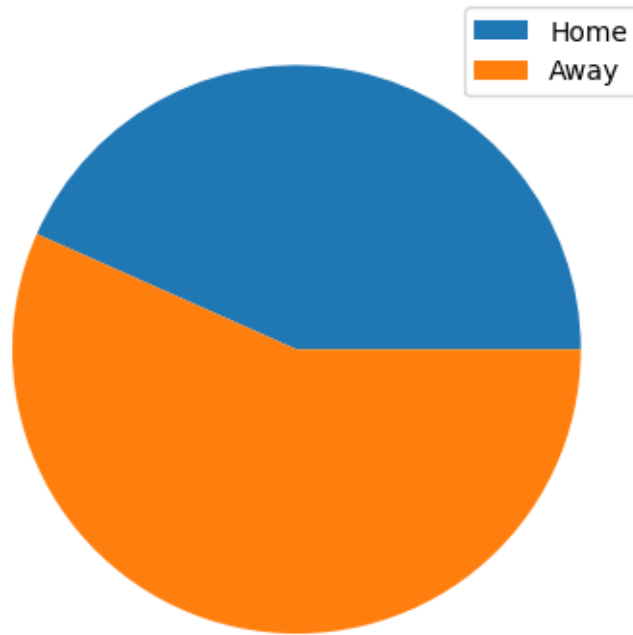
Let's get the home and away attendance for each team in 2000

```
home_away_2000 = attend_df[(attend_df["year"] == 2000) & (attend_df["week"] == 1)]  
cardinals_2000 = home_away_2000[home_away_2000["team_name"] == "Cardinals"].iloc[0]  
plt.pie([cardinals_2000["home"], cardinals_2000["away"]], labels = ["Home", "Away"])
```



Let's get the home and away attendance for each team in 2000

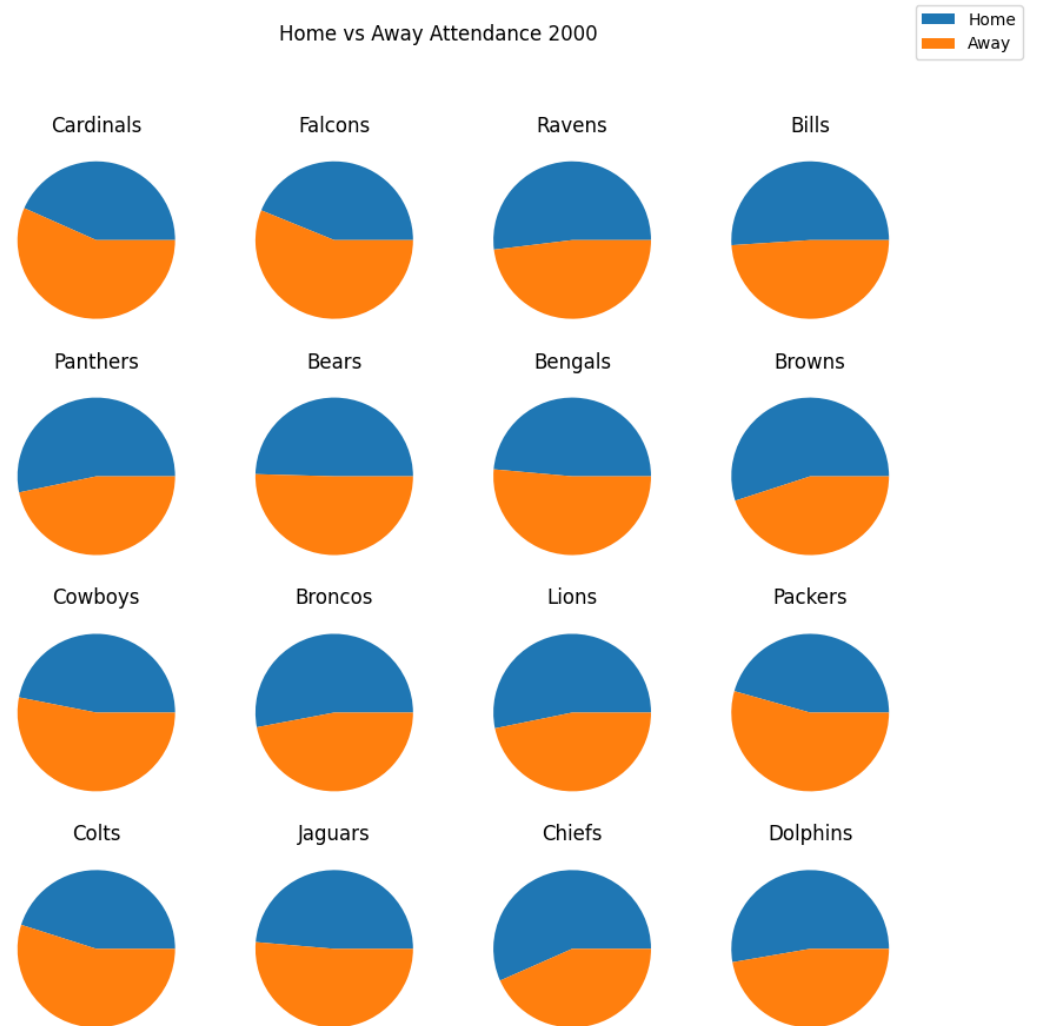
```
home_away_2000 = attend_df[(attend_df["year"] == 2000) & (attend_df["week"] == 1)]  
cardinals_2000 = home_away_2000[home_away_2000["team_name"] == "Cardinals"].iloc[0]  
plt.pie([cardinals_2000["home"], cardinals_2000["away"]])  
plt.legend(["Home", "Away"])
```




```

fig, ax = plt.subplots(4, 4, figsize=(10, 10))
for idx, row in enumerate(home_away_2000.iloc[:16].index):
    team_name = home_away_2000.loc[row, "team_name"]
    attendance = home_away_2000.loc[row, ["home", "away"]]
    ax[idx//4, idx%4].pie(attendance)
    ax[idx//4, idx%4].set_title(team_name)
fig.legend(["Home", "Away"])
fig.suptitle("Home vs Away Attendance 2000")

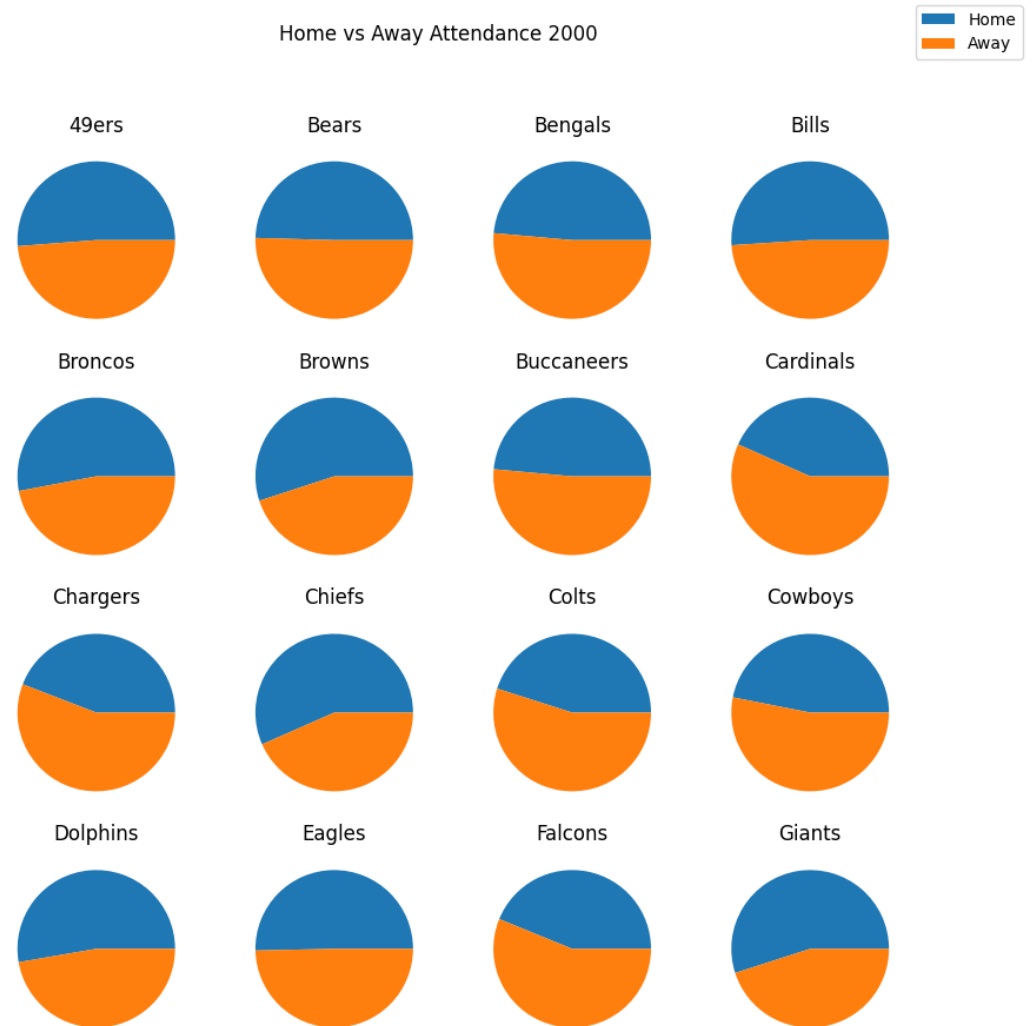
```



```

fig, ax = plt.subplots(4, 4, figsize=(10, 10))
for idx, row in enumerate(home_away_2000.sort_values("team_name").iloc[:16].index):
    team_name = home_away_2000.loc[row, "team_name"]
    attendance = home_away_2000.loc[row, ["home", "away"]]
    ax[idx//4, idx%4].pie(attendance)
    ax[idx//4, idx%4].set_title(team_name)
fig.legend(["Home", "Away"])
fig.suptitle("Home vs Away Attendance 2000")

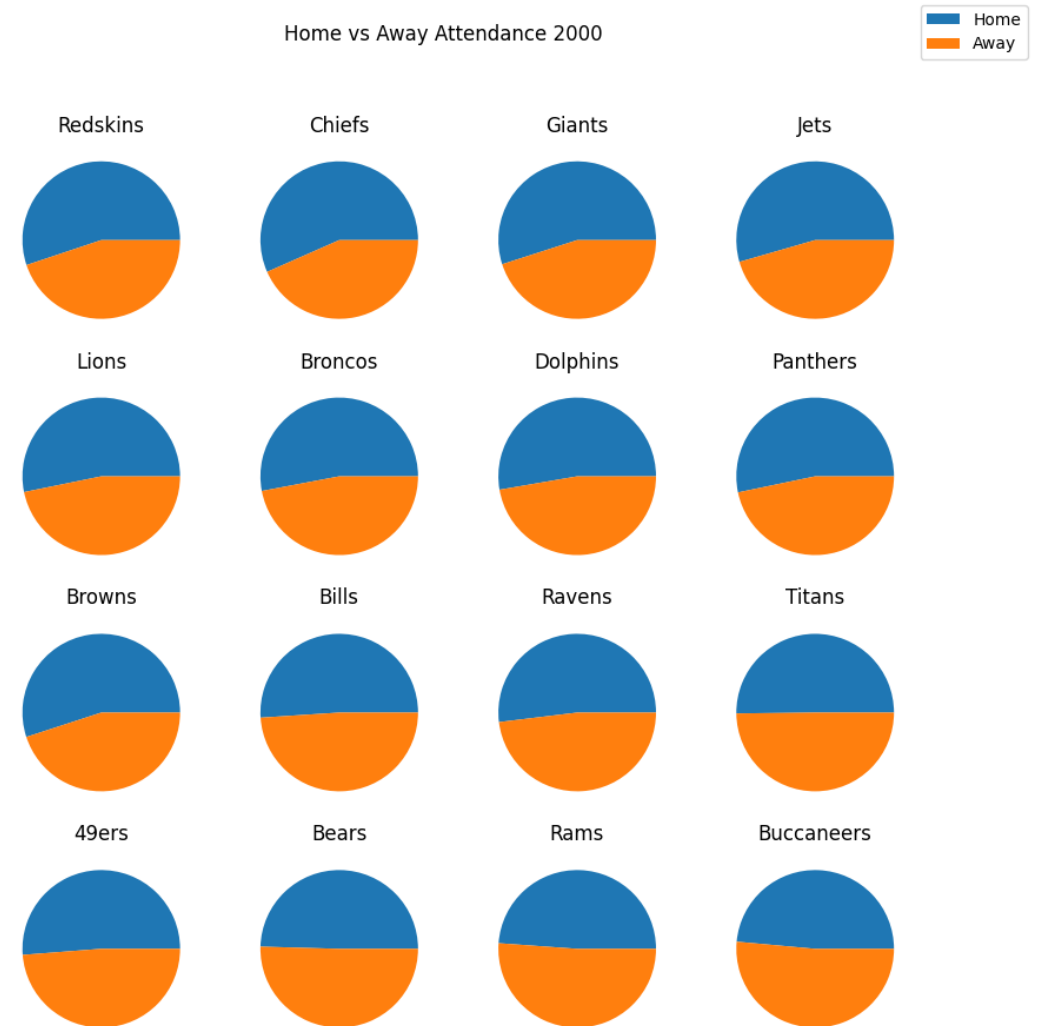
```



```

fig, ax = plt.subplots(4, 4, figsize=(10, 10))
for idx, row in enumerate(home_away_2000.sort_values("home", ascending=False).iloc[:16].index):
    team_name = home_away_2000.loc[row, "team_name"]
    attendance = home_away_2000.loc[row, ["home", "away"]]
    ax[idx//4, idx%4].pie(attendance)
    ax[idx//4, idx%4].set_title(team_name)
fig.legend(["Home", "Away"])
fig.suptitle("Home vs Away Attendance 2000")

```




```


fig, ax = plt.subplots(4, 2, figsize=(9, 14))
for idx, team_name in enumerate(["Bears", "Cowboys", "Texans", "Titans"]):
    year = 2019
    data = attendance_df[(attendance_df["year"] == year)
                        & (attendance_df["team_name"] == team_name)]
    home = data["home"].iloc[0]
    away = data["away"].iloc[0]
    ax[idx, 0].pie([home, away])
    ax[idx, 0].set_title(f"Home vs Away Attendance for {team_name}")
    ax[idx, 1].bar(np.arange(len(data)), data["weekly_attendance"])
    ax[idx, 1].set_xticks(np.arange(len(data)), data["week"])
    ax[idx, 1].set_title(f"Weekly Attendance for {team_name}")
fig.suptitle("Attendance 2019")

```

Home vs Away Attendance for Bears




Weekly Attendance for Bears




Home vs Away Attendance for Cowboys




Weekly Attendance for Cowboys




Home vs Away Attendance for Texans




Weekly Attendance for Texans




Home vs Away Attendance for Titans



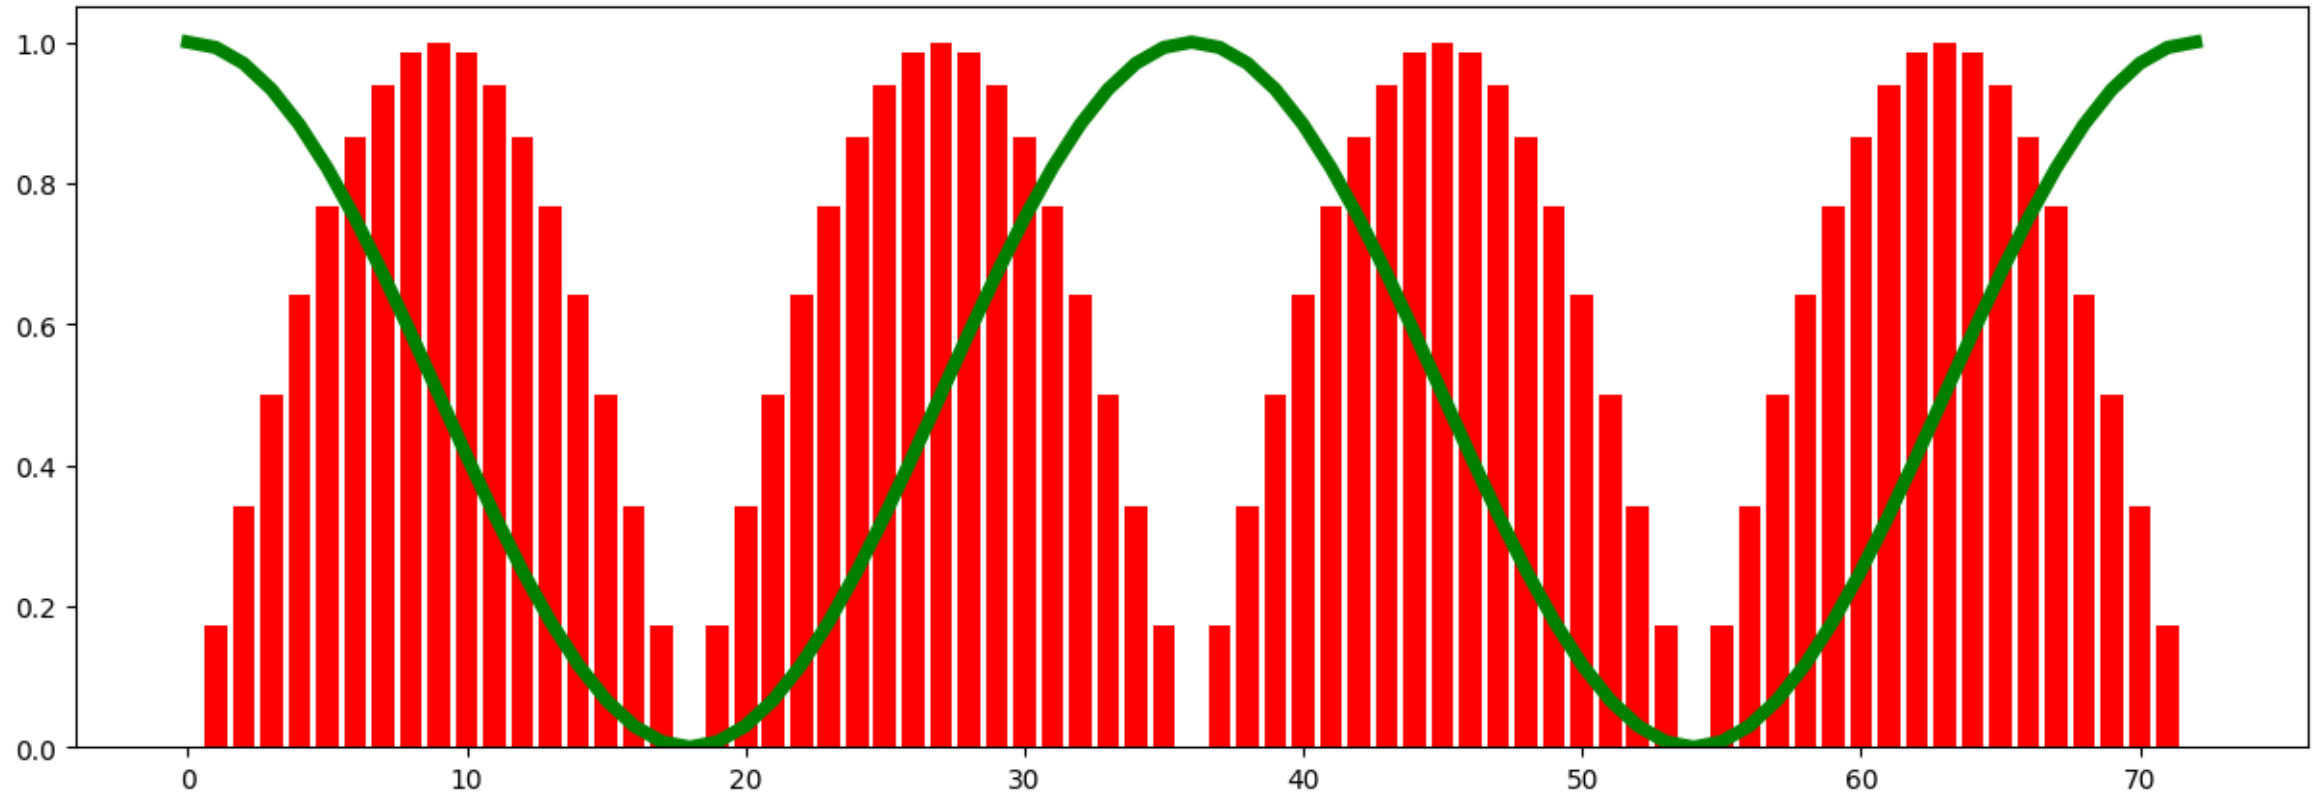
Weekly Attendance for Titans




```
x = np.linspace(0, np.pi*4, 721)
plt.plot(np.sin(x), color="k")
plt.plot(np.cos(x), linestyle="--", color="g")
plt.plot(np.sin(x) + (np.pi / 4), linestyle=":", color="c")
plt.plot(np.cos(x) - (np.pi / 4), linestyle="-.", color="r")
```




```
x = np.linspace(0, np.pi*4, 73)
plt.figure(figsize=(15,5))
plt.bar(np.arange(len(x)), np.abs(np.sin(x)), color="r")
plt.plot(np.arange(len(x)), np.cos(x) / 2 + 0.5, color="g", linewidth=5)
```




```
x = np.random.uniform(size=350)  
y = np.random.uniform(size=350)
```




```
x = np.random.uniform(size=350)
y = np.random.uniform(size=350)
c = np.random.choice(np.arange(4), size=350)
s = np.random.choice(np.arange(1, 11) * 25, size=350)
plt.scatter(x, y, c=c, s=s, alpha=0.5)
```




```
titans_attendance_2019[titans_attendance_2019["weekly_attendance"].notna()]\n    .plot(x="week", y="weekly_attendance", kind="bar")
```



```
titans_attendance_2019[titans_attendance_2019["weekly_attendance"].notna()]\n    .plot(x="week", y="weekly_attendance", kind="line")
```



Jupyter Notebook

DEMO

Three Ideas to Remember

Idea	Why	Package
Data is worthless	Big glob of numbers	numpy
Analytics are worth pennies	Discover the story of the data	pandas
Decisions are worth dollars	Tell the story of the data	matplotlib

THANK YOU!

<https://linktr.ee/douglasstarnes>